

# SO\_Tag\_Predictor\_new

November 19, 2018

```
In [2]: import warnings
        warnings.filterwarnings("ignore")
        import pandas as pd
        import sqlite3
        import csv
        import matplotlib.pyplot as plt
        import seaborn as sns
        import numpy as np
        from wordcloud import WordCloud
        import re
        import os
        from sqlalchemy import create_engine # database connection
        import datetime as dt
        from nltk.corpus import stopwords
        from nltk.tokenize import word_tokenize
        from nltk.stem.snowball import SnowballStemmer
        from sklearn.feature_extraction.text import CountVectorizer
        from sklearn.model_selection import GridSearchCV
        from sklearn.feature_extraction.text import TfidfVectorizer
        from sklearn.multiclass import OneVsRestClassifier
        from sklearn.linear_model import SGDClassifier
        from sklearn import metrics
        from sklearn.metrics import f1_score, precision_score, recall_score
        from sklearn import svm
        from sklearn.linear_model import LogisticRegression
        from skmultilearn.adapt import mlknn
        from skmultilearn.problem_transform import ClassifierChain
        from skmultilearn.problem_transform import BinaryRelevance
        from skmultilearn.problem_transform import LabelPowerset
        from sklearn.naive_bayes import GaussianNB
        from datetime import datetime
```

## 1 Stack Overflow: Tag Prediction

### 1. Business Problem

#### 1.1 Description

## Description

Stack Overflow is the largest, most trusted online community for developers to learn, share their programming knowledge, and build their careers. Stack Overflow is something which every programmer use one way or another. Each month, over 50 million developers come to Stack Overflow to learn, share their knowledge, and build their careers. It features questions and answers on a wide range of topics in computer programming. The website serves as a platform for users to ask and answer questions, and, through membership and active participation, to vote questions and answers up or down and edit questions and answers in a fashion similar to a wiki or Digg. As of April 2014 Stack Overflow has over 4,000,000 registered users, and it exceeded 10,000,000 questions in late August 2015. Based on the type of tags assigned to questions, the top eight most discussed topics on the site are: Java, JavaScript, C#, PHP, Android, jQuery, Python and HTML.

## Problem Statement

Suggest the tags based on the content that was there in the question posted on Stackoverflow.

Source: <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/>

## 1.2 Source / useful links

Data Source : <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data> Youtube : <https://youtu.be/nNDqbUhtIRg> Research paper : <https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/tagging-1.pdf>  
Research paper : <https://dl.acm.org/citation.cfm?id=2660970&dl=ACM&coll=DL>

## 1.3 Real World / Business Objectives and Constraints

1. Predict as many tags as possible with high precision and recall.
2. Incorrect tags could impact customer experience on StackOverflow.
3. No strict latency constraints.

## 2. Machine Learning problem

### 2.1 Data

#### 2.1.1 Data Overview

Refer: <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data> All of the data is in 2 files: Train and Test.

The questions are randomized and contains a mix of verbose text sites as well as sites related to math and programming. The number of questions from each site may vary, and no filtering has been performed on the questions (such as closed questions).

#### Data Field Explanation

Dataset contains 6,034,195 rows. The columns in the table are:

#### 2.1.2 Example Data point

### 2.2 Mapping the real-world problem to a Machine Learning Problem

#### 2.2.1 Type of Machine Learning Problem

It is a multi-label classification problem Multi-label Classification: Multilabel classification assigns to each sample a set of target labels. This can be thought as predicting properties of a data-point that are not mutually exclusive, such as topics that are relevant for a document. A question on Stackoverflow might be about any of C, Pointers, FileIO and/or memory-management at the same time or none of these. **Credit:** <http://scikit-learn.org/stable/modules/multiclass.html>

#### 2.2.2 Performance metric

Micro-Averaged F1-Score (Mean F Score) : The F1 score can be interpreted as a weighted average of the precision and recall, where an F1 score reaches its best value at 1 and worst score at 0. The relative contribution of precision and recall to the F1 score are equal. The formula for the F1 score is:

$F1 = 2 * (\text{precision} * \text{recall}) / (\text{precision} + \text{recall})$

In the multi-class and multi-label case, this is the weighted average of the F1 score of each class.

'Micro f1 score': Calculate metrics globally by counting the total true positives, false negatives and false positives. This is a better metric when we have class imbalance.

'Macro f1 score': Calculate metrics for each label, and find their unweighted mean. This does not take label imbalance into account.

<https://www.kaggle.com/wiki/MeanFScore> [http://scikit-learn.org/stable/modules/generated/sklearn.metrics.hamming\\_loss.html](http://scikit-learn.org/stable/modules/generated/sklearn.metrics.hamming_loss.html)  
Hamming loss : The Hamming loss is the fraction of labels that are incorrectly predicted.  
<https://www.kaggle.com/wiki/HammingLoss>

### 3. Exploratory Data Analysis

#### 3.1 Data Loading and Cleaning

##### 3.1.1 Using Pandas with SQLite to Load the data

```
In [3]: #Creating db file from csv
        #Learn SQL: https://www.w3schools.com/sql/default.asp
        if not os.path.isfile('train.db'):
            start = datetime.now()
            disk_engine = create_engine('sqlite:///train.db')
            start = dt.datetime.now()
            chunksize = 180000
            j = 0
            index_start = 1
            for df in pd.read_csv('Train.csv', names=['Id', 'Title', 'Body', 'Tags'], chunksize=
                df.index += index_start
                j+=1
                print('{} rows'.format(j*chunksize))
                df.to_sql('data', disk_engine, if_exists='append')
                index_start = df.index[-1] + 1
            print("Time taken to run this cell :", datetime.now() - start)
```

##### 3.1.2 Counting the number of rows

```
In [3]: if os.path.isfile('train.db'):
        start = datetime.now()
        con = sqlite3.connect('train.db')
        num_rows = pd.read_sql_query("""SELECT count(*) FROM data""", con)
        #Always remember to close the database
        print("Number of rows in the database :", "\n", num_rows['count(*)'].values[0])
        con.close()
        print("Time taken to count the number of rows :", datetime.now() - start)
    else:
        print("Please download the train.db file from drive or run the above cell to generate it")
```

Number of rows in the database :

6034196

Time taken to count the number of rows : 0:00:48.707495

### 3.1.3 Checking for duplicates

```
In [4]: #Learn SQL: https://www.w3schools.com/sql/default.asp
if os.path.isfile('train.db'):
    start = datetime.now()
    con = sqlite3.connect('train.db')
    df_no_dup = pd.read_sql_query('SELECT Title, Body, Tags, COUNT(*) as cnt_dup FROM data')
    con.close()
    print("Time taken to run this cell :", datetime.now() - start)
else:
    print("Please download the train.db file from drive or run the first to generate train data")
```

Time taken to run this cell : 0:02:31.241691

```
In [5]: df_no_dup.head()
# we can observe that there are duplicates
```

```
Out[5]:
```

	Title \	Body \	Tags	cnt_dup
0	Implementing Boundary Value Analysis of S...	<pre><code>#include<iosstream>\n#include<...	c++ c	1
1	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamical...	c# silverlight data-binding	1
2	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamical...	c# silverlight data-binding columns	1
3	java.lang.NoClassDefFoundError: javax/serv...	<p>I followed the guide in <a href="http://sta...	jsp jstl	1
4	java.sql.SQLException: [Microsoft][ODBC Dri...	<p>I use the following code</p>\n\n<pre><code>...	java jdbc	2

```
In [6]: print("number of duplicate questions :", num_rows['count(*)'].values[0] - df_no_dup.shape[0])
number of duplicate questions : 1827881 ( 30.292038906260256 % )
```

```
In [7]: # number of times each question appeared in our database
df_no_dup.cnt_dup.value_counts()
```

```
Out[7]: 1    2656284
        2    1272336
```

```

3      277575
4          90
5          25
6           5
Name: cnt_dup, dtype: int64

```

```

In [8]: #checking for null values
nan_rows = df_no_dup[df_no_dup.isnull().any(1)]
nan_rows

```

```

Out[8]:

```

	Title \	Body	Tags	cnt_dup
777547	Do we really need NULL?	<blockquote>\n <p><strong>Possible Duplicate:...	None	1
962680	Find all values that are not null and not in a...	<p>I am running into a problem which results i...	None	1
1126558	Handle NullObjects	<p>I have done quite a bit of research on best...	None	1
1256102	How do Germans call null	<p>In german null means 0, so how do they call...	None	1
2430668	Page cannot be null. Please ensure that this o...	<p>I get this error when i remove dynamically ...	None	1
3329908	What is the difference between NULL and "0"?	<p>What is the difference from NULL and "0"?</...>	None	1
3551595	a bit of difference between null and space	<p>I was just reading this quote</p>\n\n<block...	None	2

```

In [9]: # dropping the rows contain null value
df_no_dup.dropna(inplace=True)

```

```

In [10]: start = datetime.now()
df_no_dup["tag_count"] = df_no_dup["Tags"].apply(lambda text: len(text.split(" ")))
# adding a new feature number of tags per question
print("Time taken to run this cell :", datetime.now() - start)
df_no_dup.head()

```

Time taken to run this cell : 0:00:03.968964

```

Out[10]:

```

	Title \	Body \
0	Implementing Boundary Value Analysis of S...	<pre><code>#include<istream>\n#include&...
1	Dynamic Datagrid Binding in Silverlight?	
2	Dynamic Datagrid Binding in Silverlight?	
3	java.lang.NoClassDefFoundError: javax/serv...	
4	java.sql.SQLException: [Microsoft] [ODBC Dri...	

```

1 <p>I should do binding for datagrid dynamicall...
2 <p>I should do binding for datagrid dynamicall...
3 <p>I followed the guide in <a href="http://sta...
4 <p>I use the following code</p>\n\n<pre><code>...

```

	Tags	cnt_dup	tag_count
0	c++ c	1	2
1	c# silverlight data-binding	1	3
2	c# silverlight data-binding columns	1	4
3	jsp jstl	1	2
4	java jdbc	2	2

```

In [11]: # distribution of number of tags per question
df_no_dup.tag_count.value_counts()

```

```

Out[11]: 3    1206157
         2    1111706
         4    814996
         1    568291
         5    505158
         Name: tag_count, dtype: int64

```

```

In [12]: #Creating a new database with no duplicates
if not os.path.isfile('train_no_dup.db'):
    disk_dup = create_engine("sqlite:///train_no_dup.db")
    no_dup = pd.DataFrame(df_no_dup, columns=['Title', 'Body', 'Tags'])
    no_dup.to_sql('no_dup_train', disk_dup)

```

```

In [13]: #This method seems more appropriate to work with this much data.
#creating the connection with database file.
if os.path.isfile('train_no_dup.db'):
    start = datetime.now()
    con = sqlite3.connect('train_no_dup.db')
    tag_data = pd.read_sql_query("""SELECT Tags FROM no_dup_train""", con)
    #Always remember to close the database
    con.close()

```

```

    # Let's now drop unwanted column.
    tag_data.drop(tag_data.index[0], inplace=True)
    #Printing first 5 columns from our data frame
    tag_data.head()
    print("Time taken to run this cell :", datetime.now() - start)
else:
    print("Please download the train.db file from drive or run the above cells to genar

```

```

Time taken to run this cell : 0:00:49.989970

```

## 3.2 Analysis of Tags

### 3.2.1 Total number of unique tags

```
In [14]: # Importing & Initializing the "CountVectorizer" object, which
         # is scikit-learn's bag of words tool.

         #by default 'split()' will tokenize each tag using space.
vectorizer = CountVectorizer(tokenizer = lambda x: x.split())
# fit_transform() does two functions: First, it fits the model
# and learns the vocabulary; second, it transforms our training data
# into feature vectors. The input to fit_transform should be a list of strings.
tag_dtm = vectorizer.fit_transform(tag_data['Tags'])

In [15]: print("Number of data points :", tag_dtm.shape[0])
         print("Number of unique tags :", tag_dtm.shape[1])

Number of data points : 4206307
Number of unique tags : 42048

In [16]: #'get_feature_name()' gives us the vocabulary.
tags = vectorizer.get_feature_names()
#Lets look at the tags we have.
print("Some of the tags we have :", tags[:10])

Some of the tags we have : ['.a', '.app', '.asp.net-mvc', '.aspxauth', '.bash-profile', '.class-
```

### 3.2.3 Number of times a tag appeared

```
In [17]: # https://stackoverflow.com/questions/15115765/how-to-access-sparse-matrix-elements
         # Lets now store the document term matrix in a dictionary.
freqs = tag_dtm.sum(axis=0).A1
result = dict(zip(tags, freqs))

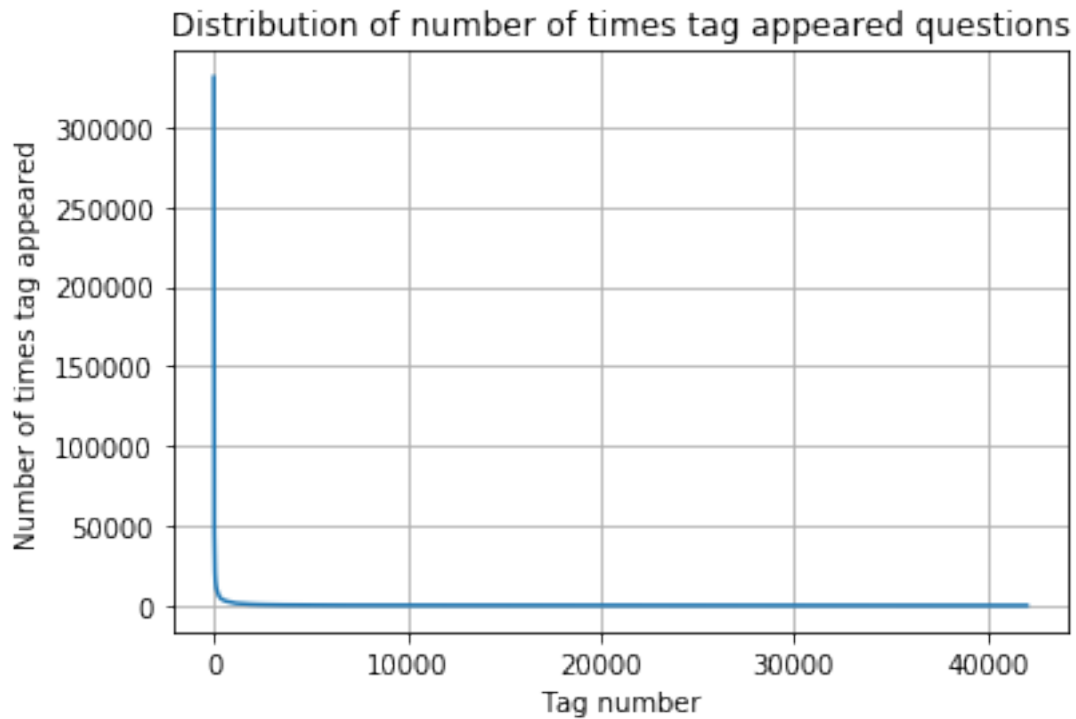
In [18]: #Saving this dictionary to csv files.
if not os.path.isfile('tag_counts_dict_dtm.csv'):
    with open('tag_counts_dict_dtm.csv', 'w') as csv_file:
        writer = csv.writer(csv_file)
        for key, value in result.items():
            writer.writerow([key, value])
tag_df = pd.read_csv("tag_counts_dict_dtm.csv", names=['Tags', 'Counts'])
tag_df.head()

Out[18]:
```

	Tags	Counts
0	jconnect	16
1	dotnetnuke-module	90
2	macromedia	22
3	ibm-jsf	8
4	rtmps	9

```
In [19]: tag_df_sorted = tag_df.sort_values(['Counts'], ascending=False)
tag_counts = tag_df_sorted['Counts'].values
```

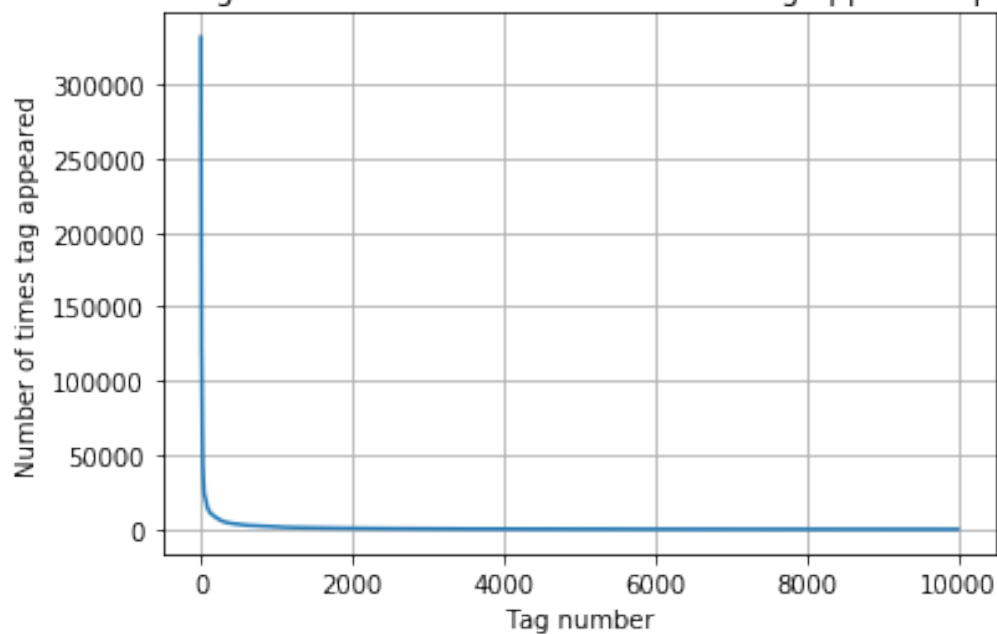
```
In [20]: plt.plot(tag_counts)
plt.title("Distribution of number of times tag appeared questions")
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
```



```
In [21]: plt.plot(tag_counts[0:10000])
plt.title('first 10k tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:10000:25]), tag_counts[0:10000:25])
```



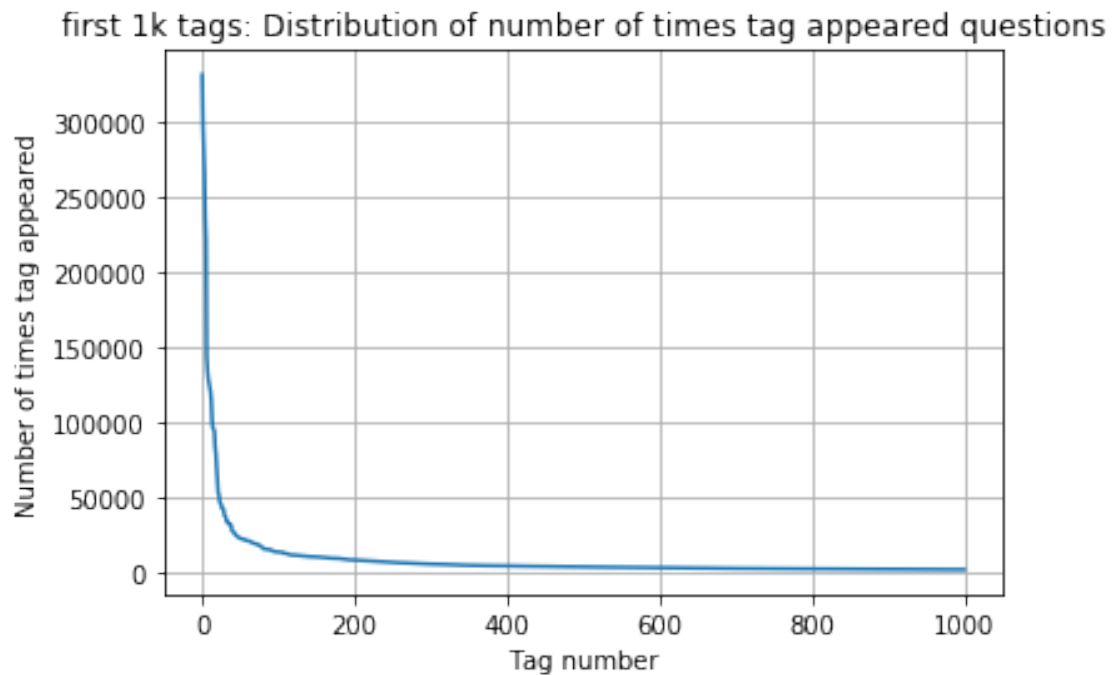
first 10k tags: Distribution of number of times tag appeared questions



400	[331505	44829	22429	17728	13364	11162	10029	9148	8054	7151
6466	5865	5370	4983	4526	4281	4144	3929	3750	3593	
3453	3299	3123	2986	2891	2738	2647	2527	2431	2331	
2259	2186	2097	2020	1959	1900	1828	1770	1723	1673	
1631	1574	1532	1479	1448	1406	1365	1328	1300	1266	
1245	1222	1197	1181	1158	1139	1121	1101	1076	1056	
1038	1023	1006	983	966	952	938	926	911	891	
882	869	856	841	830	816	804	789	779	770	
752	743	733	725	712	702	688	678	671	658	
650	643	634	627	616	607	598	589	583	577	
568	559	552	545	540	533	526	518	512	506	
500	495	490	485	480	477	469	465	457	450	
447	442	437	432	426	422	418	413	408	403	
398	393	388	385	381	378	374	370	367	365	
361	357	354	350	347	344	342	339	336	332	
330	326	323	319	315	312	309	307	304	301	
299	296	293	291	289	286	284	281	278	276	
275	272	270	268	265	262	260	258	256	254	
252	250	249	247	245	243	241	239	238	236	
234	233	232	230	228	226	224	222	220	219	
217	215	214	212	210	209	207	205	204	203	
201	200	199	198	196	194	193	192	191	189	
188	186	185	183	182	181	180	179	178	177	
175	174	172	171	170	169	168	167	166	165	
164	162	161	160	159	158	157	156	156	155	

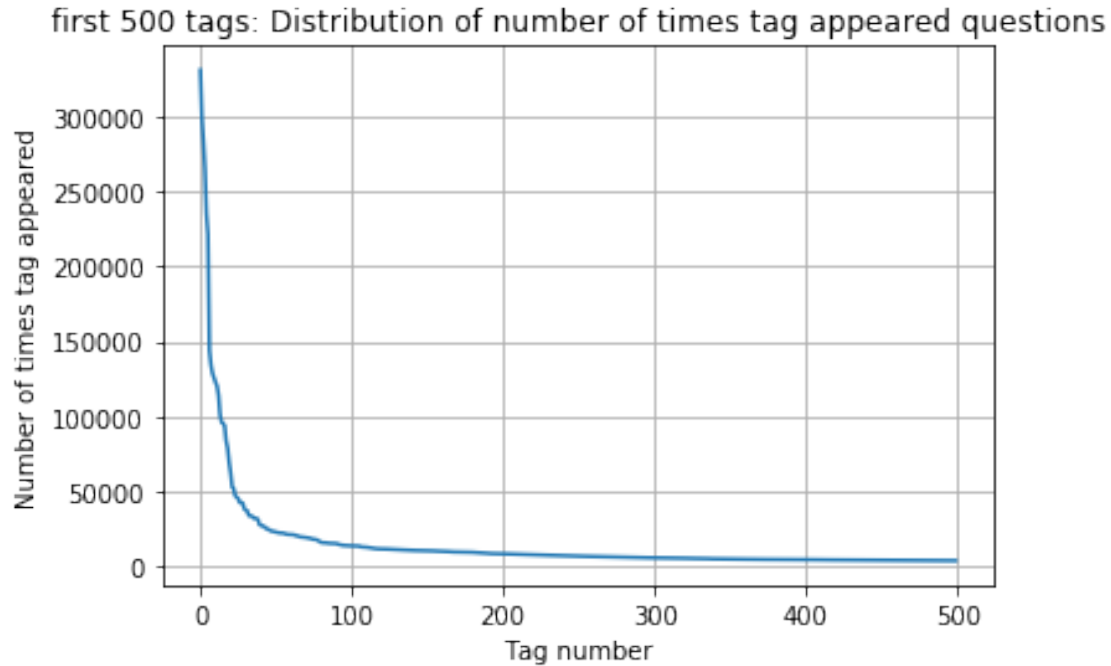
154	153	152	151	150	149	149	148	147	146
145	144	143	142	142	141	140	139	138	137
137	136	135	134	134	133	132	131	130	130
129	128	128	127	126	126	125	124	124	123
123	122	122	121	120	120	119	118	118	117
117	116	116	115	115	114	113	113	112	111
111	110	109	109	108	108	107	106	106	106
105	105	104	104	103	103	102	102	101	101
100	100	99	99	98	98	97	97	96	96
95	95	94	94	93	93	93	92	92	91
91	90	90	89	89	88	88	87	87	86
86	86	85	85	84	84	83	83	83	82
82	82	81	81	80	80	80	79	79	78
78	78	78	77	77	76	76	76	75	75
75	74	74	74	73	73	73	73	72	72]

```
In [22]: plt.plot(tag_counts[0:1000])
plt.title('first 1k tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:1000:5]), tag_counts[0:1000:5])
```



200	[331505	221533	122769	95160	62023	44829	37170	31897	26925	24537
22429	21820	20957	19758	18905	17728	15533	15097	14884	13703	
13364	13157	12407	11658	11228	11162	10863	10600	10350	10224	
10029	9884	9719	9411	9252	9148	9040	8617	8361	8163	
8054	7867	7702	7564	7274	7151	7052	6847	6656	6553	
6466	6291	6183	6093	5971	5865	5760	5577	5490	5411	
5370	5283	5207	5107	5066	4983	4891	4785	4658	4549	
4526	4487	4429	4335	4310	4281	4239	4228	4195	4159	
4144	4088	4050	4002	3957	3929	3874	3849	3818	3797	
3750	3703	3685	3658	3615	3593	3564	3521	3505	3483	
3453	3427	3396	3363	3326	3299	3272	3232	3196	3168	
3123	3094	3073	3050	3012	2986	2983	2953	2934	2903	
2891	2844	2819	2784	2754	2738	2726	2708	2681	2669	
2647	2621	2604	2594	2556	2527	2510	2482	2460	2444	
2431	2409	2395	2380	2363	2331	2312	2297	2290	2281	
2259	2246	2222	2211	2198	2186	2162	2142	2132	2107	
2097	2078	2057	2045	2036	2020	2011	1994	1971	1965	
1959	1952	1940	1932	1912	1900	1879	1865	1855	1841	
1828	1821	1813	1801	1782	1770	1760	1747	1741	1734	
1723	1707	1697	1688	1683	1673	1665	1656	1646	1639]	

```
In [23]: plt.plot(tag_counts[0:500])
plt.title('first 500 tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:500:5]), tag_counts[0:500:5])
```



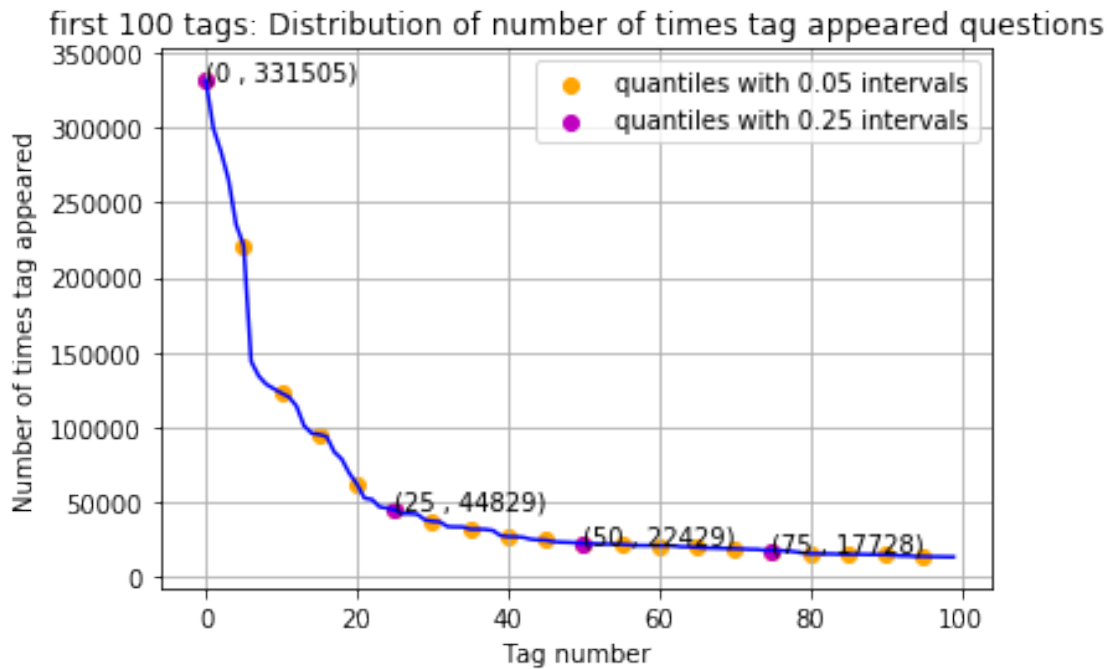
100	331505	221533	122769	95160	62023	44829	37170	31897	26925	24537
22429	21820	20957	19758	18905	17728	15533	15097	14884	13703	
13364	13157	12407	11658	11228	11162	10863	10600	10350	10224	
10029	9884	9719	9411	9252	9148	9040	8617	8361	8163	
8054	7867	7702	7564	7274	7151	7052	6847	6656	6553	
6466	6291	6183	6093	5971	5865	5760	5577	5490	5411	
5370	5283	5207	5107	5066	4983	4891	4785	4658	4549	
4526	4487	4429	4335	4310	4281	4239	4228	4195	4159	
4144	4088	4050	4002	3957	3929	3874	3849	3818	3797	
3750	3703	3685	3658	3615	3593	3564	3521	3505	3483	

```
In [24]: plt.plot(tag_counts[0:100], c='b')
plt.scatter(x=list(range(0,100,5)), y=tag_counts[0:100:5], c='orange', label="quantiles
# quantiles with 0.25 difference
plt.scatter(x=list(range(0,100,25)), y=tag_counts[0:100:25], c='m', label = "quantiles

for x,y in zip(list(range(0,100,25)), tag_counts[0:100:25]):
    plt.annotate(text="({} , {})".format(x,y), xy=(x,y), xytext=(x-0.05, y+500))

plt.title('first 100 tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.legend()
```

```
plt.show()
print(len(tag_counts[0:100:5]), tag_counts[0:100:5])
```



```
20 [331505 221533 122769 95160 62023 44829 37170 31897 26925 24537
    22429 21820 20957 19758 18905 17728 15533 15097 14884 13703]
```

```
In [25]: # Store tags greater than 10K in one list
lst_tags_gt_10k = tag_df[tag_df.Counts>10000].Tags
#Print the length of the list
print ('{} Tags are used more than 10000 times'.format(len(lst_tags_gt_10k)))
# Store tags greater than 100K in one list
lst_tags_gt_100k = tag_df[tag_df.Counts>100000].Tags
#Print the length of the list.
print ('{} Tags are used more than 100000 times'.format(len(lst_tags_gt_100k)))
```

```
153 Tags are used more than 10000 times
14 Tags are used more than 100000 times
```

Observations: 1. There are total 153 tags which are used more than 10000 times. 2. 14 tags are used more than 100000 times. 3. Most frequent tag (i.e. c#) is used 331505 times. 4. Since some tags occur much more frequently than others, Micro-averaged F1-score is the appropriate metric for this problem.

### 3.2.4 Tags Per Question

```
In [26]: #Storing the count of tag in each question in list 'tag_count'
tag_quest_count = tag_dtm.sum(axis=1).tolist()
#Converting each value in the 'tag_quest_count' to integer.
tag_quest_count=[int(j) for i in tag_quest_count for j in i]
print ('We have total {} datapoints.'.format(len(tag_quest_count)))

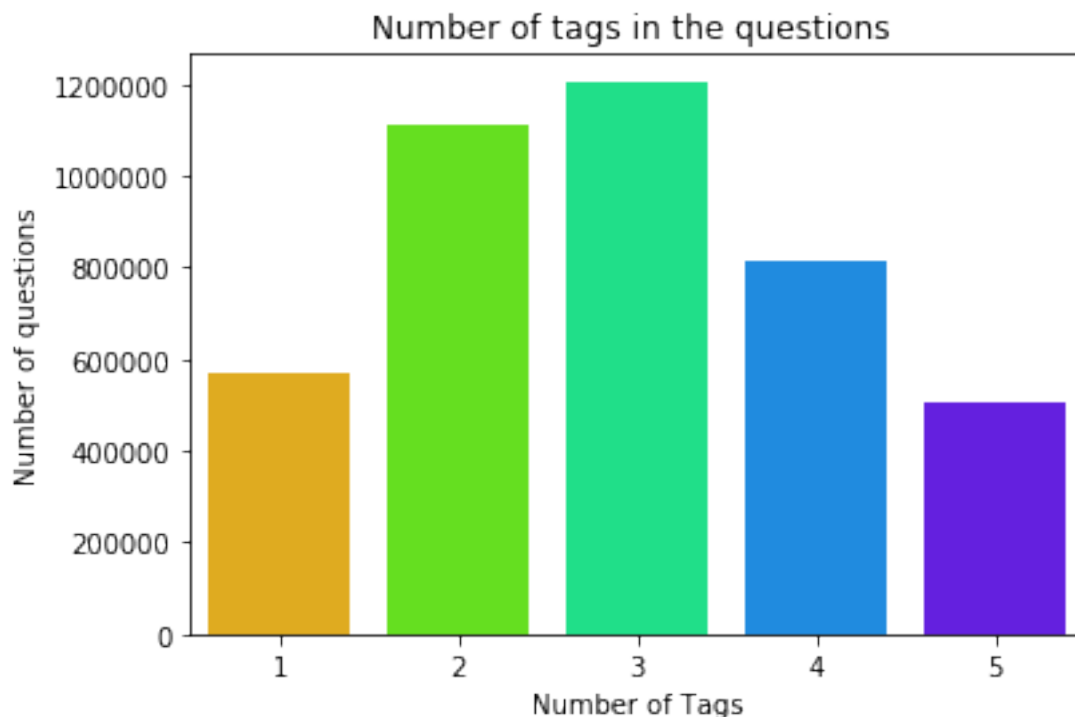
print(tag_quest_count[:5])
```

We have total 4206307 datapoints.  
[3, 4, 2, 2, 3]

```
In [27]: print( "Maximum number of tags per question: %d"%max(tag_quest_count))
print( "Minimum number of tags per question: %d"%min(tag_quest_count))
print( "Avg. number of tags per question: %f"% ((sum(tag_quest_count)*1.0)/len(tag_quest_count)))
```

Maximum number of tags per question: 5  
Minimum number of tags per question: 1  
Avg. number of tags per question: 2.899443

```
In [28]: sns.countplot(tag_quest_count, palette='gist_rainbow')
plt.title("Number of tags in the questions ")
plt.xlabel("Number of Tags")
plt.ylabel("Number of questions")
plt.show()
```

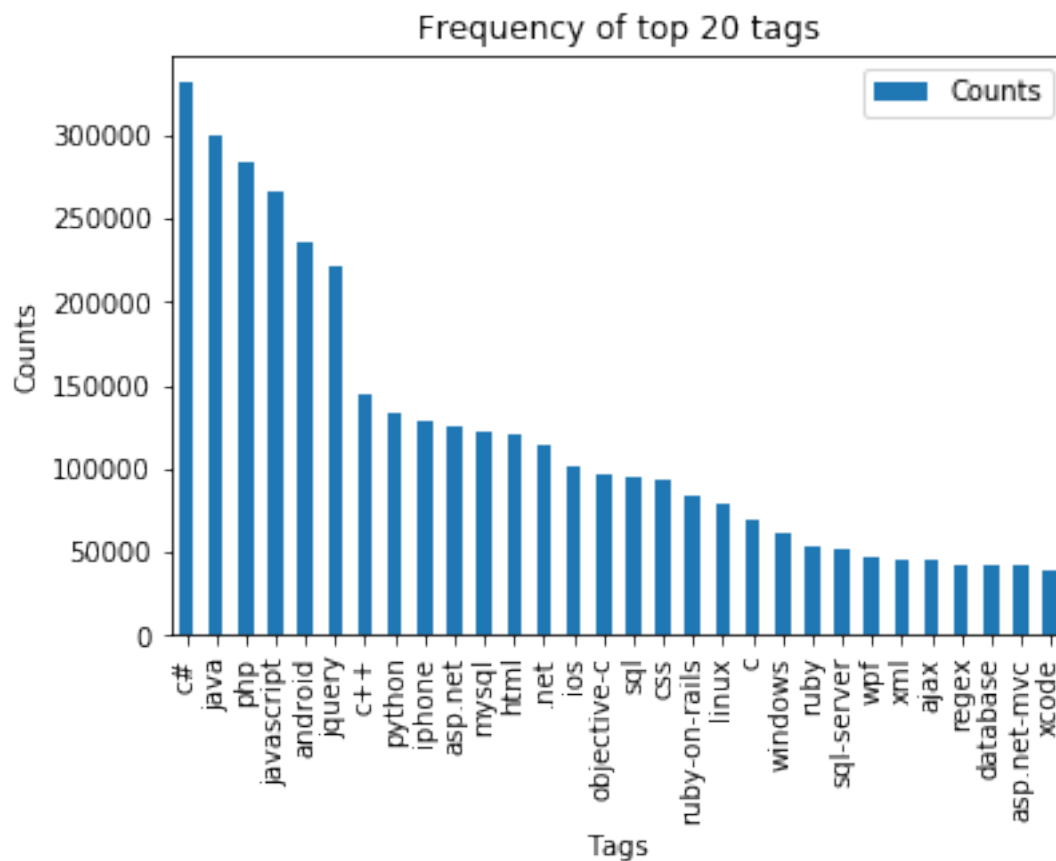




Observations: A look at the word cloud shows that "c#", "java", "php", "asp.net", "javascript", "c++" are some of the most frequent tags.

### 3.2.6 The top 20 tags

```
In [30]: i=np.arange(30)
tag_df_sorted.head(30).plot(kind='bar')
plt.title('Frequency of top 20 tags')
plt.xticks(i, tag_df_sorted['Tags'])
plt.xlabel('Tags')
plt.ylabel('Counts')
plt.show()
```



Observations: 1. Majority of the most frequent tags are programming language. 2. C# is the top most frequent programming language. 3. Android, IOS, Linux and windows are among the top most frequent operating systems.

### 3.3 Cleaning and preprocessing of Questions

```
In [3]: def striphtml(data):
cleanr = re.compile('<.*?>')
cleantext = re.sub(cleanr, ' ', str(data))
return cleantext
```



```
stop_words = set(stopwords.words('english'))
stemmer = SnowballStemmer("english")
```

```
In [4]: #http://www.sqlitetutorial.net/sqlite-python/create-tables/
def create_connection(db_file):
    """ create a database connection to the SQLite database
        specified by db_file
    :param db_file: database file
    :return: Connection object or None
    """
    try:
        conn = sqlite3.connect(db_file)
        return conn
    except Error as e:
        print(e)

    return None

def create_table(conn, create_table_sql):
    """ create a table from the create_table_sql statement
    :param conn: Connection object
    :param create_table_sql: a CREATE TABLE statement
    :return:
    """
    try:
        c = conn.cursor()
        c.execute(create_table_sql)
    except Error as e:
        print(e)

def checkTableExists(dbcon):
    cursr = dbcon.cursor()
    str = "select name from sqlite_master where type='table'"
    table_names = cursr.execute(str)
    print("Tables in the databse:")
    tables = table_names.fetchall()
    print(tables[0][0])
    return(len(tables))

def create_database_table(database, query):
    conn = create_connection(database)
    if conn is not None:
        create_table(conn, query)
        checkTableExists(conn)
    else:
        print("Error! cannot create the database connection.")
    conn.close()
```

```
sql_create_table = """CREATE TABLE IF NOT EXISTS QuestionsProcessed (question text NOT N
create_database_table("Processed.db", sql_create_table)
```

Tables in the database:  
QuestionsProcessed

\_\_ We will sample the number of tags instead considering all of them (due to limitation of computing power) \_\_

```
In [5]: def tags_to_choose(n):
        t = multilabel_y.sum(axis=0).tolist()[0]
        sorted_tags_i = sorted(range(len(t)), key=lambda i: t[i], reverse=True)
        multilabel_yn=multilabel_y[:,sorted_tags_i[:n]]
        return multilabel_yn

        def questions_explained_fn(n):
            multilabel_yn = tags_to_choose(n)
            x= multilabel_yn.sum(axis=1)
            return (np.count_nonzero(x==0))
```

```
In [6]: sql_create_table = """CREATE TABLE IF NOT EXISTS QuestionsProcessed (question text NOT N
        create_database_table("Titlemoreweight.db", sql_create_table)
```

Tables in the database:  
QuestionsProcessed

```
In [7]: # http://www.sqlitetutorial.net/sqlite-delete/
        # https://stackoverflow.com/questions/2279706/select-random-row-from-a-sqlite-table

        read_db = 'train_no_dup.db'
        write_db = 'Titlemoreweight.db'
        train_datasize = 400000
        if os.path.isfile(read_db):
            conn_r = create_connection(read_db)
            if conn_r is not None:
                reader =conn_r.cursor()
                # for selecting first 0.5M rows
                reader.execute("SELECT Title, Body, Tags From no_dup_train LIMIT 500001;")
                # for selecting random points
                #reader.execute("SELECT Title, Body, Tags From no_dup_train ORDER BY RANDOM() LI

        if os.path.isfile(write_db):
            conn_w = create_connection(write_db)
            if conn_w is not None:
                tables = checkTableExists(conn_w)
                writer =conn_w.cursor()
                if tables != 0:
```

```
writer.execute("DELETE FROM QuestionsProcessed WHERE 1")
print("Cleared All the rows")
```

Tables in the database:

QuestionsProcessed

Cleared All the rows

### 3.3.1 Preprocessing of questions

```
<li> Sample 0.5M data points and taking just 500 most important tags </li>
<li> Separate Code from Body </li>
<li> Remove Special characters from Question title and description (not in code)</li>
<li> <b> Give more weightage to title : Add title three times to the question </b> </li>

<li> Remove stop words (Except 'C') </li>
<li> Remove HTML Tags </li>
<li> Convert all the characters into small letters </li>
<li> Use SnowballStemmer to stem the words </li>
```

```
In [8]: #http://www.bernzilla.com/2008/05/13/selecting-a-random-row-from-an-sqlite-table/
start = datetime.now()
preprocessed_data_list=[]
reader.fetchone()
questions_with_code=0
len_pre=0
len_post=0
questions_processed = 0
for row in reader:

    is_code = 0

    title, question, tags = row[0], row[1], str(row[2])

    if '<code>' in question:
        questions_with_code+=1
        is_code = 1
    x = len(question)+len(title)
    len_pre+=x

    code = str(re.findall(r'<code>(.*?)</code>', question, flags=re.DOTALL))

    question=re.sub('<code>(.*?)</code>', '', question, flags=re.MULTILINE|re.DOTALL)
    question=striphtml(question.encode('utf-8'))

    title=title.encode('utf-8')

    # adding title three time to the data to increase its weight
    # add tags string to the training data
```

```

question=str(title)+" "+str(title)+" "+str(title)+" "+question

#     if questions_proccesed<=train_datasize:
#         question=str(title)+" "+str(title)+" "+str(title)+" "+question+" "+str(tags)
#     else:
#         question=str(title)+" "+str(title)+" "+str(title)+" "+question

question=re.sub(r'[^A-Za-z0-9#+.\-]+',' ',question)
words=word_tokenize(str(question.lower()))

#Removing all single letter and and stopwords from question exceptt for the letter '
question=' '.join(str(stemmer.stem(j)) for j in words if j not in stop_words and (len(j)>1))

len_post+=len(question)
tup = (question,code,tags,x,len(question),is_code)
questions_proccesed += 1
writer.execute("insert into QuestionsProcessed(question,code,tags,words_pre,words_post) values(%s,%s,%s,%s,%s)"%tup)
if (questions_proccesed%100000==0):
    print("number of questions completed=",questions_proccesed)

no_dup_avg_len_pre=(len_pre*1.0)/questions_proccesed
no_dup_avg_len_post=(len_post*1.0)/questions_proccesed

print( "Avg. length of questions(Title+Body) before processing: %d"%no_dup_avg_len_pre)
print( "Avg. length of questions(Title+Body) after processing: %d"%no_dup_avg_len_post)
print( "Percent of questions containing code: %d"%((questions_with_code*100.0)/questions_proccesed))

print("Time taken to run this cell :", datetime.now() - start)

number of questions completed= 100000
number of questions completed= 200000
number of questions completed= 300000
number of questions completed= 400000
number of questions completed= 500000
Avg. length of questions(Title+Body) before processing: 1239
Avg. length of questions(Title+Body) after processing: 424
Percent of questions containing code: 57
Time taken to run this cell : 0:17:07.249072

In [9]: # never forget to close the conections or else we will end up with database locks
conn_r.commit()
conn_w.commit()
conn_r.close()
conn_w.close()

```

\_\_ Sample quesitons after preprocessing of data \_\_

```
In [10]: if os.path.isfile(write_db):
        conn_r = create_connection(write_db)
        if conn_r is not None:
            reader = conn_r.cursor()
            reader.execute("SELECT question From QuestionsProcessed LIMIT 10")
            print("Questions after preprocessed")
            print('='*100)
            reader.fetchone()
            for row in reader:
                print(row)
                print('-'*100)
        conn_r.commit()
        conn_r.close()
```

Questions after preprocessed

```
=====
('dynam datagrid bind silverlight dynam datagrid bind silverlight dynam datagrid bind silverlight
-----
('java.lang.noclassdeffounderror javax servlet jsp tagext taglibraryvalid java.lang.noclassdeffo
-----
('java.sql.sqllexcept microsoft odbc driver manag invalid descriptor index java.sql.sqllexcept mic
-----
('better way updat feed fb php sdk better way updat feed fb php sdk better way updat feed fb php
-----
('btnadd click event open two window record ad btnadd click event open two window record ad btna
-----
('sql inject issu prevent correct form submiss php sql inject issu prevent correct form submiss
-----
('countabl subaddit lebesgu measur countabl subaddit lebesgu measur countabl subaddit lebesgu me
-----
('hql equival sql queri hql equival sql queri hql equival sql queri hql queri replac name class
-----
('undefin symbol architectur i386 objc class skpsmtpmessag referenc error undefin symbol archite
-----
```

\_\_ Saving Preprocessed data to a Database \_\_

```
In [11]: #Taking 0.5 Million entries to a dataframe.
        write_db = 'Titlemoreweight.db'
        if os.path.isfile(write_db):
            conn_r = create_connection(write_db)
            if conn_r is not None:
                preprocessed_data = pd.read_sql_query("""SELECT question, Tags FROM QuestionsPr
            conn_r.commit()
            conn_r.close()
```

```
In [12]: preprocessed_data.head()
```

```

Out [12]:                                     question \
0 dynam datagrid bind silverlight dynam datagrid...
1 dynam datagrid bind silverlight dynam datagrid...
2 java.lang.noclassdeffoundererror javax servlet j...
3 java.sql.sqlexcept microsoft odbc driver manag...
4 better way updat feed fb php sdk better way up...

                                     tags
0          c# silverlight data-binding
1 c# silverlight data-binding columns
2                                     jsp jstl
3                                     java jdbc
4          facebook api facebook-php-sdk

```

```

In [13]: print("number of data points in sample :", preprocessed_data.shape[0])
         print("number of dimensions :", preprocessed_data.shape[1])

```

```

number of data points in sample : 500000
number of dimensions : 2

```

\_\_ Converting string Tags to multilable output variables \_\_

```

In [14]: vectorizer = CountVectorizer(tokenizer = lambda x: x.split(), binary='true')
         multilabel_y = vectorizer.fit_transform(preprocessed_data['tags'])

```

\_\_ Selecting 500 Tags \_\_

```

In [15]: questions_explained = []
         total_tags=multilabel_y.shape[1]
         total_qs=preprocessed_data.shape[0]
         for i in range(500, total_tags, 100):
             questions_explained.append(np.round(((total_qs-questions_explained_fn(i))/total_qs)

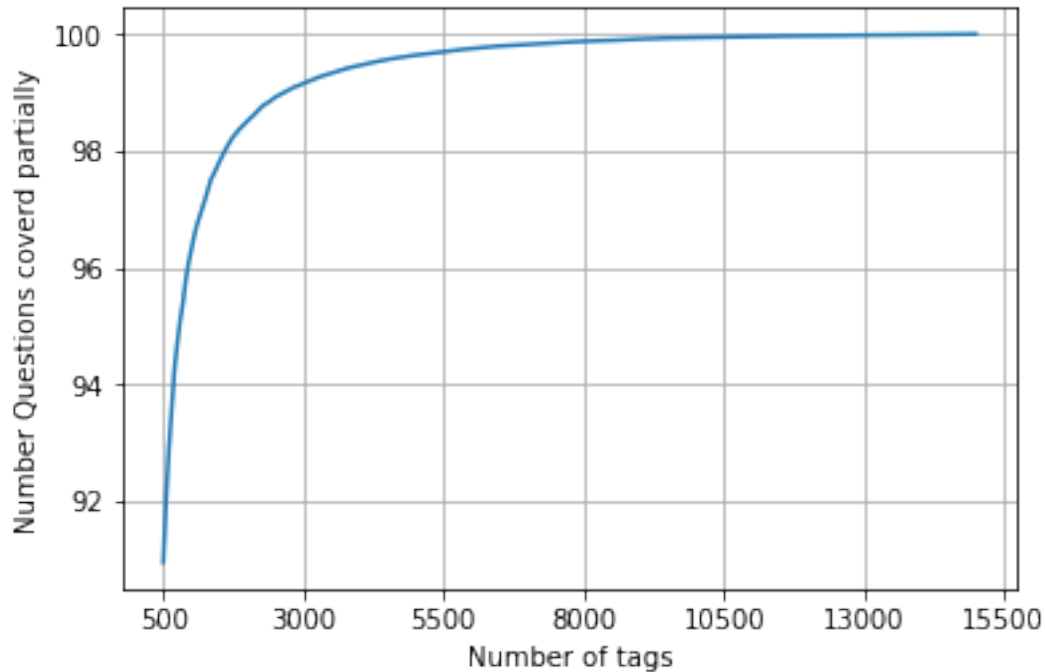
```

```

In [16]: fig, ax = plt.subplots()
         ax.plot(questions_explained)
         xlabel = list(500+np.array(range(-50,450,50))*50)
         ax.set_xticklabels(xlabel)
         plt.xlabel("Number of tags")
         plt.ylabel("Number Questions coverd partially")
         plt.grid()
         plt.show()

         # you can choose any number of tags based on your computing power, minimun is 500(it ca
         print("with ",5500,"tags we are covering ",questions_explained[50],"% of questions")
         print("with ",500,"tags we are covering ",questions_explained[0],"% of questions")

```



with 5500 tags we are covering 99.157 % of questions  
 with 500 tags we are covering 90.956 % of questions

```
In [17]: # we will be taking 500 tags
         multilabel_yx = tags_to_choose(500)
         print("number of questions that are not covered :", questions_explained_fn(500),"out of 500000")

number of questions that are not covered : 45221 out of 500000
```

```
In [18]: x_train=preprocessed_data.head(train_datasize)
         x_test=preprocessed_data.tail(preprocessed_data.shape[0] - 400000)

         y_train = multilabel_yx[0:train_datasize,:]
         y_test = multilabel_yx[train_datasize:preprocessed_data.shape[0],:]
```

```
In [19]: print("Number of data points in train data :", y_train.shape)
         print("Number of data points in test data :", y_test.shape)
```

Number of data points in train data : (400000, 500)  
 Number of data points in test data : (100000, 500)

#### 4. Modeling.

## 4.1 Modeling using Tfidf vectorizer

### 4.1.1 Featurizing data with Tfidf vectorizer

```
In [20]: start = datetime.now()
         vectorizer = TfidfVectorizer(min_df=0.00009, max_features=200000, smooth_idf=True, norm='l2',
                                     tokenizer = lambda x: x.split(), sublinear_tf=False, ngram_range=(1,1))
         x_train_multilabel = vectorizer.fit_transform(x_train['question'])
         x_test_multilabel = vectorizer.transform(x_test['question'])
         print("Time taken to run this cell :", datetime.now() - start)
```

Time taken to run this cell : 0:05:06.846641

```
In [21]: print("Dimensions of train data X:",x_train_multilabel.shape, "Y :",y_train.shape)
         print("Dimensions of test data X:",x_test_multilabel.shape,"Y:",y_test.shape)
```

Dimensions of train data X: (400000, 94927) Y : (400000, 500)

Dimensions of test data X: (100000, 94927) Y: (100000, 500)

## 4.1.2 Logistic Regression with OneVsRestClassifier using Tfidf vectorizer

### 4.1.2.1 Hyperparameter tuning

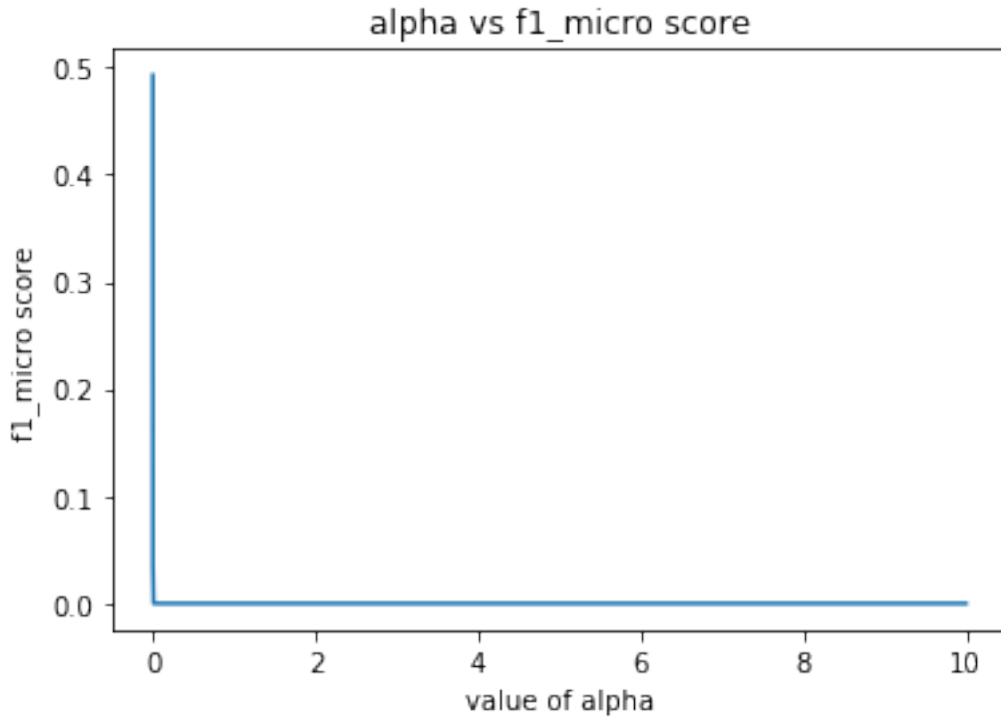
```
In [25]: param={'estimator__alpha': [10**-5, 10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1]}
         classifier = OneVsRestClassifier(SGDClassifier(loss='log', penalty='l1'))
         gsv = GridSearchCV(estimator = classifier, param_grid=param, cv=3, verbose=0, scoring='f1_micro')
         gsv.fit(x_train_multilabel, y_train)
```

```
best_alpha = gsv.best_estimator_.get_params()['estimator__alpha']
print('value of alpha after hyperparameter tuning : ',best_alpha)
print('-----')
# plotting C vs f1_micro_score
x_1=[]
y_1=[]
for x in gsv.grid_scores_:
    x_1.append(x[0]['estimator__alpha'])
    y_1.append(x[1])
plt.plot(x_1,y_1)
plt.xlabel('value of alpha')
plt.ylabel('f1_micro score')
plt.title('alpha vs f1_micro score')
plt.show()
```

value of alpha after hyperparameter tuning : 1e-05

-----





#### 4.1.2.2 Applying model using best hyperparameter

```
In [26]: start = datetime.now()
         #best_alpha = gsv.best_estimator_.get_params()['estimator__alpha']
         classifier = OneVsRestClassifier(SGDClassifier(loss='log', alpha=best_alpha, penalty='l1'))
         classifier.fit(x_train_multilabel, y_train)
         predictions = classifier.predict (x_test_multilabel)

         print("Accuracy :",metrics.accuracy_score(y_test, predictions))
         print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

         precision = precision_score(y_test, predictions, average='micro')
         recall = recall_score(y_test, predictions, average='micro')
         f1 = f1_score(y_test, predictions, average='micro')

         print("Micro-average quality numbers")
         print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

         precision = precision_score(y_test, predictions, average='macro')
         recall = recall_score(y_test, predictions, average='macro')
         f1 = f1_score(y_test, predictions, average='macro')

         print("Macro-average quality numbers")
```

```

print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall,
# print (metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.23644
Hamming loss  0.00278178
Micro-average quality numbers
Precision: 0.7211, Recall: 0.3258, F1-measure: 0.4488
Macro-average quality numbers
Precision: 0.5478, Recall: 0.2573, F1-measure: 0.3340
Time taken to run this cell : 0:05:02.703501

```

### 4.1.3 Linear SVM with OneVsRestClassifier using Tfidf vectorizer

#### 4.1.3.1 Hyperparameter tuning

```

In [28]: param={'estimator__alpha': [10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1]}
classifier = OneVsRestClassifier(SGDClassifier(loss='hinge', penalty='l1'))
gsv = GridSearchCV(estimator = classifier, param_grid=param, cv=3, verbose=0, scoring='
gsv.fit(x_train_multilabel, y_train)

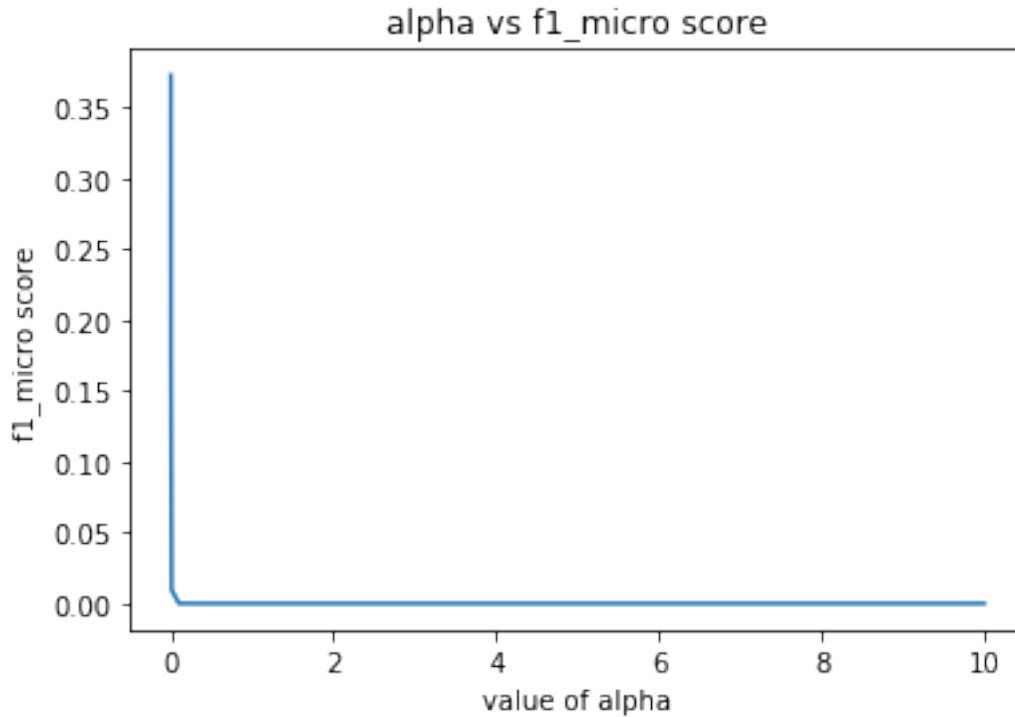
best_alpha = gsv.best_estimator_.get_params()['estimator__alpha']
print('value of alpha after hyperparameter tuning : ', best_alpha)
print('-----')
# plotting C vs f1_micro_score
x_1=[]
y_1=[]
for x in gsv.grid_scores_:
    x_1.append(x[0]['estimator__alpha'])
    y_1.append(x[1])
plt.plot(x_1, y_1)
plt.xlabel('value of alpha')
plt.ylabel('f1_micro score')
plt.title('alpha vs f1_micro score')
plt.show()

```

```

value of alpha after hyperparameter tuning : 0.0001
-----

```



#### 4.1.3.2 Applying model using best hyperparameter

```
In [29]: start = datetime.now()
         #best_alpha = gsv.best_estimator_.get_params()['estimator__alpha']
         classifier = OneVsRestClassifier(SGDClassifier(loss='hinge', alpha=best_alpha, penalty=
         classifier.fit(x_train_multilabel, y_train)
         predictions = classifier.predict (x_test_multilabel)

         print("Accuracy :",metrics.accuracy_score(y_test, predictions))
         print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

         precision = precision_score(y_test, predictions, average='micro')
         recall = recall_score(y_test, predictions, average='micro')
         f1 = f1_score(y_test, predictions, average='micro')

         print("Micro-average quality numbers")
         print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

         precision = precision_score(y_test, predictions, average='macro')
         recall = recall_score(y_test, predictions, average='macro')
         f1 = f1_score(y_test, predictions, average='macro')

         print("Macro-average quality numbers")
```

```

print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall,
# print (metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.2105
Hamming loss  0.0029036
Micro-average quality numbers
Precision: 0.8169, Recall: 0.2123, F1-measure: 0.3370
Macro-average quality numbers
Precision: 0.2440, Recall: 0.1298, F1-measure: 0.1607
Time taken to run this cell : 0:05:01.558172

```

## 4.2 Modeling using Count vectorizer

### 4.2.1 Featurizing data with Count vectorizer

```

In [21]: start = datetime.now()
         vectorizer = CountVectorizer(min_df=0.00009, max_features=200000, \
                                     tokenizer = lambda x: x.split(), ngram_range=(1,4))
         x_train_multilabel = vectorizer.fit_transform(x_train['question'])
         x_test_multilabel = vectorizer.transform(x_test['question'])
         print("Time taken to run this cell :", datetime.now() - start)

```

```

Time taken to run this cell : 0:11:02.103345

```

```

In [22]: print("Dimensions of train data X:",x_train_multilabel.shape, "Y :",y_train.shape)
         print("Dimensions of test data X:",x_test_multilabel.shape,"Y:",y_test.shape)

```

```

Dimensions of train data X: (400000, 95585) Y : (400000, 500)
Dimensions of test data X: (100000, 95585) Y: (100000, 500)

```

## 4.2.2 Logistic Regression with OneVsRestClassifier using count vectorizer

### 4.2.2.1 Hyperparameter tuning

```

In [56]: param={'estimator__alpha': [10**-5, 10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1]}
         classifier = OneVsRestClassifier(SGDClassifier(loss='log', penalty='l1'))
         gsv = GridSearchCV(estimator = classifier, param_grid=param, cv=3, verbose=0, scoring='
         gsv.fit(x_train_multilabel, y_train)

         param={'estimator__alpha': [10**-5, 10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1]}
         classifier = OneVsRestClassifier(SGDClassifier(loss='log', penalty='l1'))
         gsv = GridSearchCV(estimator = classifier, param_grid=param, cv=3, verbose=0, scoring='
         gsv.fit(x_train_multilabel, y_train)

         best_alpha = gsv.best_estimator_.get_params()['estimator__alpha']
         print('value of alpha after hyperparameter tuning : ',best_alpha)

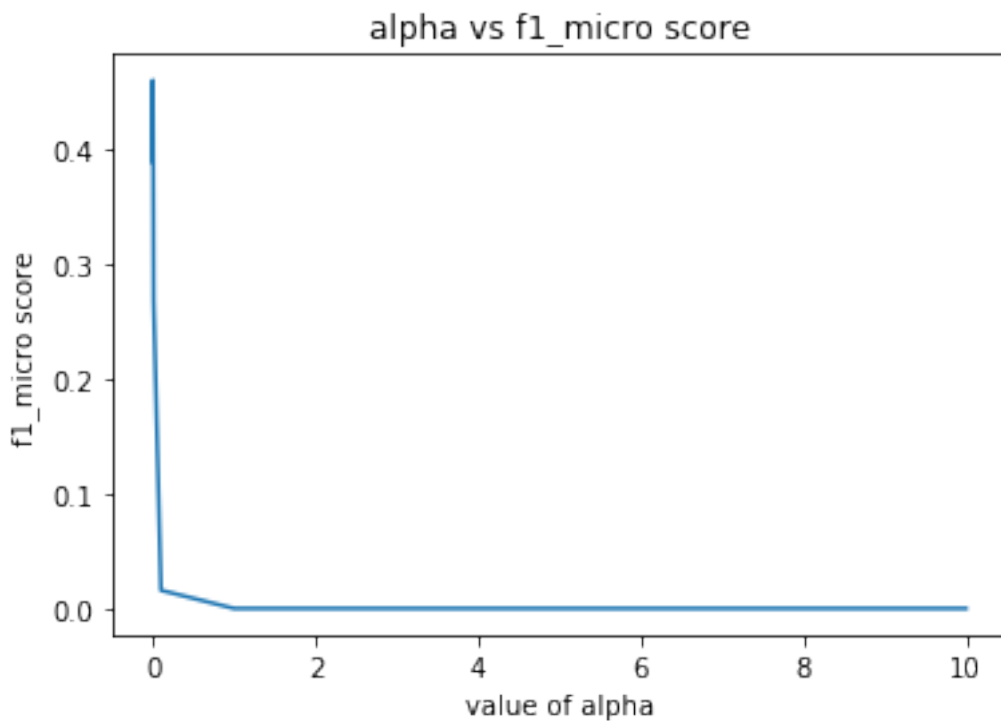
```

```

print('-----')
# plotting C vs f1_micro_score
x_1=[]
y_1=[]
for x in gsv.grid_scores_:
    x_1.append(x[0]['estimator__alpha'])
    y_1.append(x[1])
plt.plot(x_1,y_1)
plt.xlabel('value of alpha')
plt.ylabel('f1_micro score')
plt.title('alpha vs f1_micro score')
plt.show()

```

value of alpha after hyperparameter tuning : 0.001



#### 4.2.2.2 Applying model using best hyperparameter

```

In [57]: start = datetime.now()
         #best_alpha = gsv.best_estimator_.get_params()['estimator__alpha']
         classifier = OneVsRestClassifier(SGDClassifier(loss='log', alpha=best_alpha, penalty='l1'))
         classifier.fit(x_train_multilabel, y_train)
         predictions = classifier.predict (x_test_multilabel)

```

```

print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall,

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall,

#print (metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.18621
Hamming loss  0.00322218
Micro-average quality numbers
Precision: 0.5636, Recall: 0.3238, F1-measure: 0.4113
Macro-average quality numbers
Precision: 0.4073, Recall: 0.2397, F1-measure: 0.2823

Time taken to run this cell : 0:05:26.013286

```

## 4.2.3 Linear SVM with OneVsRestClassifier

### 4.2.3.1 Hyperparameter tuning

```

In [25]: param={'estimator__alpha': [10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1]}
         classifier = OneVsRestClassifier(SGDClassifier(loss='hinge', penalty='l1'))
         gsv = GridSearchCV(estimator = classifier, param_grid=param, cv=3, verbose=0, scoring='
         gsv.fit(x_train_multilabel, y_train)

param={'estimator__alpha': [10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1]}
classifier = OneVsRestClassifier(SGDClassifier(loss='hinge', penalty='l1'))
gsv = GridSearchCV(estimator = classifier, param_grid=param, cv=3, verbose=0, scoring='
gsv.fit(x_train_multilabel, y_train)

best_alpha = gsv.best_estimator_.get_params()['estimator__alpha']
print('value of alpha after hyperparameter tuning : ',best_alpha)
print('-----')

```

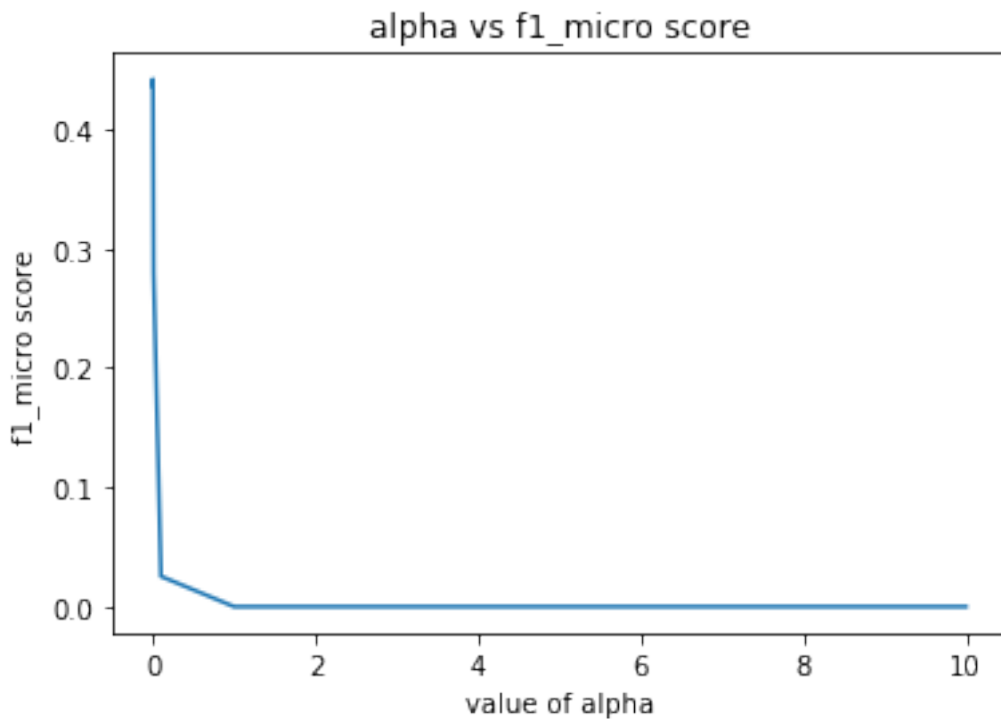
```

# plotting C vs f1_micro_score
x_1=[]
y_1=[]
for x in gsv.grid_scores_:
    x_1.append(x[0]['estimator__alpha'])
    y_1.append(x[1])
plt.plot(x_1,y_1)
plt.xlabel('value of alpha')
plt.ylabel('f1_micro score')
plt.title('alpha vs f1_micro score')
plt.show()

```

value of alpha after hyperparameter tuning : 0.001

---



#### 4.2.3.2 Applying model using best hyperparameter

```

In [26]: start = datetime.now()
         #best_alpha = gsv.best_estimator_.get_params()['estimator__alpha']
         classifier = OneVsRestClassifier(SGDClassifier(loss='hinge', alpha=best_alpha, penalty=
         classifier.fit(x_train_multilabel, y_train)
         predictions = classifier.predict (x_test_multilabel)

```

```

print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall,

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall,

#print (metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.17942
Hamming loss 0.00326302
Micro-average quality numbers
Precision: 0.5525, Recall: 0.3226, F1-measure: 0.4074
Macro-average quality numbers
Precision: 0.3128, Recall: 0.2396, F1-measure: 0.2549
Time taken to run this cell : 0:05:42.001801

```

## 1.1 Performance Table

Sr. No.	Model	Featurization	Micro f1_score	Macro f1_score	Hamming loss	Accuracy
1	Logistic Regression	Tfidf vectorizer	0.4488	0.3340	0.0027	0.2364
2	Linear SVM	Tfidf vectorizer	0.3370	0.1607	0.0029	0.2105
3	Logistic Regression	Count vectorizer	0.4113	0.2823	0.0032	0.1862
4	Linear SVM	Count vectorizer	0.4074	0.2549	0.0032	0.1794

## 1.2 Conclusion

- We have choosen 'f1\_micro' scoring metric because of the stated business statement.



- Used bag of words upto 4 grams and Tfidf upto 3 grams.
- For logistic regression , I have used 'SGDClassifier' instead of 'LogisticRegression'. The reason is 'LogisticRegression' takes lots of time for hyperparameter tuning. Even we have not choosen any complex model like xgboost, because the dimension is very high and linear model works fairly well in high dimension and the complex model like xgboost may not work well for this much high dimension, as well as it takes lots of time for hyperparameter tuning.
- We can see in the performance table that Logistic Regression with Tfidf vectorizer works better than Linear SVM.